

LINMA2710 - Scientific Computing

Single Instruction Multiple Data (SIMD)

P.-A. Absil and B. Legat

☐ Full Width Mode ☐ Present Mode

Table of Contents

Motivation

SIMD inspection


Auto-Vectorization

Interleaving

Fused instructions

Motivation ⇄


The need for parallelism ⇄



[Image source](#)

A bit of historical context ↻

- **1972** : C language created by Dennis Ritchie and Ken Thompson to ease development of Unix (previously developed in **assembly**)
- **1985** : C++ created by Bjarne Stroustrup
- **2003** : LLVM started at University of Illinois
- **2005** : Apple hires Chris Lattner from the university
- **2007** : He then creates the LLVM-based compiler Clang
- **2009** : Mozilla start developing an LLVM-based compiler for Rust
- **2009** : Development starts on Julia, with LLVM-based compiler



A sum function in C and Julia ↻

```
float sum(float *vec, int length) {  
    float total = 0;  
    for (int i = 0; i < length; i++) {  
        total += vec[i];  
    }  
    return total;  
}
```

```
1 c_sum(x::Vector{Cfloat}) = ccall(("sum", sum_float_lib), Cfloat, (Ptr{Cfloat},  
    Cint), x, length(x));
```

julia_sum (generic function with 1 method)

```
1 function julia_sum(v::Vector{T}) where {T}  
2     total = zero(T)  
3     for i in eachindex(v)  
4         total += v[i]  
5     end  
6     return total  
7 end
```

Let's make a small benchmark

```
vec_float =
```

```
► [0.152564, 0.943411, 0.806974, 0.617942, 0.770227, 0.454559, 0.832874, 0.194885, 0.668397
```

```
1 vec_float = rand(Float32, 2^16)
```

```
32775.977f0
```

```
1 @btime c_sum($vec_float)
```



```
242.514 μs (0 allocations: 0 bytes)
```



```
32775.977f0
```

```
1 @btime julia_sum($vec_float)
```



```
60.633 μs (0 allocations: 0 bytes)
```



► How to speed up the C code ?


⚠ PlutoRunner.embed_display(x) is deprecated and will be removed in a future release. Use AbstractPlutoDingetjes.Display.@embed(x) instead.

Tip

As accessing global variables is slow in Julia, it is important to add `$` in front of them when using `@btime`. This is less critical in Pluto though as it handles global variables differently. To see why, try removing the `$`, you should see `1` allocations instead of zero.

```
float sum(float *vec, int length) {  
    float total = 0;  
    for (int i = 0; i < length; i++) {  
        total += vec[i];  
    }  
    return total;  
}
```

Summing with SIMD $\Leftrightarrow$



Faster Julia code $\Leftrightarrow$

► How to get the same speed up from the Julia code ?

⚠ PlutoRunner.embed_display(x) is deprecated and will be removed in a future release. Use AbstractPlutoDingetjes.Display.@embed(x) instead.

```
julia_sum_fast (generic function with 1 method)
1 function julia_sum_fast(v::Vector{T}) where {T}
2     total = zero(T)
3     for i in eachindex(v)
4         @fastmath total += @inbounds v[i]
5     end
6     return total
7 end
```

32775.996f0

```
1 @btime julia_sum_fast($vec_float)
```

2.544 μs (0 allocations: 0 bytes)

julia_sum_simd (generic function with 1 method)

```
1 function julia_sum_simd(v::Vector{T}) where {T}
2     total = zero(T)
3     @simd for i in eachindex(v)
4         total += v[i]
5     end
6     return total
7 end
```

32775.996f0

```
1 @btime julia_sum_simd($vec_float)
```

2.544 μs (0 allocations: 0 bytes)

Careful with fast math ↗

► Why are the three elements in the center of the vector ignored in this example ?

⚠ PlutoRunner.embed_display(x) is deprecated and will be removed in a future release. Use AbstractPlutoDingetjes.Display.@embed(x) instead.

```
test_kahan = ▶ [1.0, 2.98023f-8, 2.98023f-8, 2.98023f-8, 0.000119209]
```

```
1 test_kahan = Cfloat[1.0, eps(Cfloat)/4, eps(Cfloat)/4, eps(Cfloat)/4,
1000eps(Cfloat)]
```

1.000119298696518

```
1 sum(Float64.(test_kahan))
```

1.0001192f0

```
1 c_sum(test_kahan[[1, 5]])
```

1.0001192f0

```
1 c_sum(test_kahan)
```

To improve the accuracy this, we consider the [Kahan summation algorithm](#).

```
1.0001193f0
```

```
1 c_sum_kahan(test_kahan)
```

Optimization level : ▼

Enable -ffast-math ? ☐

► **What happens when -ffast-math is enabled ?**

⚠ `PlutoRunner.embed_display(x)` is deprecated and will be removed in a future release. Use `AbstractPlutoDingetjes.Display.@embed(x)` instead.

For further details, see [this blog post](#).

Tip

`eps` gives the difference between `1` and the number closest to `1`. See also `prevfloat` and `nextfloat`.

```
float sum_kahan(float* vec, int length) {
    float total, c, t, y;
    int i;
    total = c = 0.0f;
    for (i = 0; i < length; i++) {
        y = vec[i] - c;
        t = total + y;
        c = (t - total) - y;
        total = t;
    }
    return total;
}
```

SIMD inspection ↗

Instruction sets ↗

The data is **packed** on a single SIMD unit whose width and register depends on the instruction set family. The single instruction is then run in parallel on all elements of this small **vector** stored in the SIMD unit. These give the prefix `vp` to the instruction names that stands from *Vectorized Packed*.

Instruction Set Family	Width of SIMD unit	Register
Streaming SIMD Extension (SSE)	128-bit	%xmm
Advanced Vector Extensions (AVX)	256-bit	%ymm
AVX-512	512-bit	%zmm

```
► ProcessChain([Process(`lscpu`, ProcessExited(0)), Process(`grep Flag`, ProcessExited(0))
```

```
1 run(pipeline(`lscpu`, `grep Flag`))
```

⌵

Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2 ht syscall nx m mxext fxsr_opt pdpe1gb rdtscp lm constant_tsc rep_good noopl tsc_reliable nonst op_tsc cpuid extd_apicid aperfmperf tsc_known_freq pni pclmulqdq ssse3 fma cx1 6 pcid sse4_1 sse4_2 movbe popcnt aes xsave avx f16c rdrand hypervisor lahf_lm cmp_legacy svm cr8_legacy abm sse4a misalignsse 3dnowprefetch osvw topoext vmm call fsgsbase bmi1 avx2 smep bmi2 erms invpcid rdseed adx smap clflushopt clwb sha_ni xsaveopt xsavec xgetbv1 xsaves user_shstk clzero xsaveerptr rdpru arat npt nrrip_save tsc_scale vmcb_clean flushbyasid decodeassists pausefilter pfthr eshold v_vmsave_vmload umip vaes vpcmlmulqdq rdpid fsrm

Tip

To determine which instruction set is supported for your computer, look at the `Flags` list in the output of `lscpu`. We can check in the [Intel® Intrinsic Guide](#) that `avx`, `avx2` and `avx_vnni` are in the AVX family.

SIMD at LLVM level

How can you check that SIMD is enabled? Let's check at the level of LLVM IR.

f (generic function with 1 method)

```
1 function f(x1, x2, x3, x4, y1, y2, y3, y4)
2     z1 = x1 + y1
3     z2 = x2 + y2
4     z3 = x3 + y3
5     z4 = x4 + y4
6     return z1, z2, z3, z4
7 end
```

```
1 @code_llvm debuginfo=:none f(1, 2, 3, 4, 5, 6, 7, 8)
```



```
; Function Signature: f(Int64, Int64, Int64, Int64, Int64, Int64, Int64, Int64)
define void @julia_f_34855(ptr noalias nocapture noundef nonnull sret([4 x i64]) align 8 dereferenceable(32) %sret_return, i64 signext %"x1::Int64", i64 signext %"x2::Int64", i64 signext %"x3::Int64", i64 signext %"x4::Int64", i64 signext %"y1::Int64", i64 signext %"y2::Int64", i64 signext %"y3::Int64", i64 signext %"y4::Int64") #0 {
top:
    %0 = add i64 %"y1::Int64", %"x1::Int64"
    %1 = add i64 %"y2::Int64", %"x2::Int64"
    %2 = add i64 %"y3::Int64", %"x3::Int64"
    %3 = add i64 %"y4::Int64", %"x4::Int64"
    store i64 %0, ptr %sret_return, align 8
    %"new::Tuple.sroa.2.0.sret_return.sroa_idx" = getelementptr inbounds i8, ptr %sret_return, i64 8
    store i64 %1, ptr %"new::Tuple.sroa.2.0.sret_return.sroa_idx", align 8
    %"new::Tuple.sroa.3.0.sret_return.sroa_idx" = getelementptr inbounds i8, ptr %sret_return, i64 16
    store i64 %2, ptr %"new::Tuple.sroa.3.0.sret_return.sroa_idx", align 8
    %"new::Tuple.sroa.4.0.sret_return.sroa_idx" = getelementptr inbounds i8, ptr %sret_return, i64 24
    store i64 %3, ptr %"new::Tuple.sroa.4.0.sret_return.sroa_idx", align 8
    ret void
}
```

Tip

If we see `add i64`, it means that each `Int64` is added independently

Packing the data to enable SIMD

f_broadcast (generic function with 1 method)

```
1 function f_broadcast(x, y)
2     z = x .+ y
3     return z
4 end
```

```
1 @code_llvm debuginfo=:none f_broadcast((1, 2, 3, 4), (1, 2, 3, 4))
```

```
> ; Function Signature: f_broadcast(NTuple{4, Int64}, NTuple{4, Int64})
define void @julia_f_broadcast_33503(ptr noalias nocapture noundef nonnull sret([4 x i64]) align 8 dereferenceable(32) %sret_return, ptr nocapture noundef nonnull readonly align 8 dereferenceable(32) %"x::Tuple", ptr nocapture noundef nonnull readonly align 8 dereferenceable(32) %"y::Tuple") #0 {
top:
    %0 = load <4 x i64>, ptr %"x::Tuple", align 8
    %1 = load <4 x i64>, ptr %"y::Tuple", align 8
    %2 = add <4 x i64> %1, %0
    store <4 x i64> %2, ptr %sret_return, align 8
    ret void
}
```

Tip

`load <4 x i64>` means that 4 `Int64` are loaded into a 256-bit wide SIMD unit.

SIMD at assembly level ↗

```
1 @code_native debuginfo=:none f_broadcast((1, 2, 3, 4), (1, 2, 3, 4))
```

```
.text
.file "f_broadcast"
.section .ltext,"axl",@progbits
.globl julia_f_broadcast_33807 # -- Begin function julia_f_broadc
ast_33807
.p2align 4, 0x90
.type julia_f_broadcast_33807,@function
julia_f_broadcast_33807: # @julia_f_broadcast_33807
; Function Signature: f_broadcast(NTuple{4, Int64}, NTuple{4, Int64})
# %bb.0: # %top
#DEBUG_VALUE: f_broadcast:x <- [$rsi+0]
#DEBUG_VALUE: f_broadcast:y <- [$rdx+0]
push rbp
vmovdqu ymm0, ymmword ptr [rdx]
mov rbp, rsp
mov rax, rdi
vpaddq ymm0, ymm0, ymmword ptr [rsi]
vmovdqu ymmword ptr [rdi], ymm0
pop rbp
vzeroupper
ret
.Lfunc_end0:
.size julia_f_broadcast_33807, .Lfunc_end0-julia_f_broadcast_33807
# -- End function
.type ".L+Core.Tuple#33809",@object # @".L+Core.Tuple#33809"
.section .lrodata,"al",@progbits
.p2align 3, 0x0
".L+Core.Tuple#33809":
quad ".L+Core.Tuple#33809_ii"
```

Tip

The suffix **v** in front of the instruction stands for **vectorized**. It means it is using a SIMD unit.

Tuples implementing the array interface ↗

N =  2

```

1 let
2     T = Float64
3     A = rand(SMatrix{N,N,T})
4     x = rand(SVector{N,T})
5     @code_llvm debuginfo=:none A * x
6 end

```



```

; Function Signature: *(StaticArraysCore.SArray{Tuple{2, 2}, Float64, 2,
4}, StaticArraysCore.SArray{Tuple{2}, Float64, 1, 2})
define void @"julia_*)_34800"(ptr noalias nocapture noundef nonnull sret([1 x
[2 x double]]) align 8 dereferenceable(16) %sret_return, ptr nocapture noundef
nonnull readonly align 8 dereferenceable(32) %"A::SArray", ptr nocapture noundef
nonnull readonly align 8 dereferenceable(16) %"B::SArray") #0 {
top:
    %"A::SArray[3]_ptr" = getelementptr inbounds i8, ptr %"A::SArray", i64 16
    %0 = load <2 x double>, ptr %"B::SArray", align 8
    %1 = load <2 x double>, ptr %"A::SArray", align 8
    %2 = shufflevector <2 x double> %0, <2 x double> poison, <2 x i32> zeroiniti
alizer
    %3 = fmul <2 x double> %1, %2
    %4 = load <2 x double>, ptr %"A::SArray[3]_ptr", align 8
    %5 = shufflevector <2 x double> %0, <2 x double> poison, <2 x i32> <i32 1, i
32 1>
    %6 = fmul contract <2 x double> %4, %5
    %7 = fadd contract <2 x double> %3, %6
    store <2 x double> %7, ptr %sret_return, align 8
    ret void
}

```

Tip

Small arrays that are allocated on the stack like tuples and implemented in `StaticArrays.jl`. Operating on them leverages SIMD.

Auto-Vectorization ⇄

LLVM Loop Vectorizer for a C array ⇄

```
1 emit_llvm(c_sum_code_for_llvm, cflags = [sum_opt; sum_flags]);
```

```
; ModuleID = '/tmp/jl_ZOCdT9/main.c'
source_filename = "/tmp/jl_ZOCdT9/main.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-f8
0:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

; Function Attrs: noinline nounwind optnone uwtable
define dso_local i32 @sum(ptr noundef %0, i32 noundef %1) #0 {
    %3 = alloca ptr, align 8
    %4 = alloca i32, align 4
    %5 = alloca i32, align 4
    %6 = alloca i32, align 4
    store ptr %0, ptr %3, align 8
    store i32 %1, ptr %4, align 4
    store i32 0, ptr %5, align 4
    store i32 0, ptr %6, align 4
    br label %7

7:                                     ; preds = %19, %2
    %8 = load i32, ptr %6, align 4
    %9 = load i32, ptr %4, align 4
    %10 = icmp slt i32 %8, %9
    br i1 %10, label %11, label %22

11:                                   ; preds = %7
    %12 = load ptr, ptr %3, align 8
    %13 = load i32, ptr %6, align 4
    %14 = sext i32 %13 to i64
    %15 = getelementptr inbounds i32, ptr %12, i64 %14
```

```
int sum(int *vec, int length) {
    int total = 0;
    for (int i = 0; i < length; i++) {
        total += vec[i];
    }
    return total;
}
```

No pragma ▼

No pragma ▼

No pragma ▼

Element type :

Optimization level :

☐ -msse3

☐ -mavx2

☐ -mavx512f

☐ -ffast-math

☐ -march=native

⚠ PlutoRunner.embed_display(x) is deprecated and will be removed in a future release. Use AbstractPlutoDingetjes.Display.@embed(x) instead.

LLVM Loop Vectorizer for a C++ vector ↗

```
> ; ModuleID = '/tmp/jl_UfRUC8/main.c'
source_filename = "/tmp/jl_UfRUC8/main.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

; Function Attrs: noinline nounwind optnone uwtable
define dso_local i32 @sum(ptr noundef %0, i32 noundef %1) #0 {
    %3 = alloca ptr, align 8
    %4 = alloca i32, align 4
    %5 = alloca i32, align 4
    %6 = alloca i32, align 4
    store ptr %0, ptr %3, align 8
    store i32 %1, ptr %4, align 4
    store i32 0, ptr %5, align 4
    store i32 0, ptr %6, align 4
    br label %7

7:                                     ; preds = %19, %2
    %8 = load i32, ptr %6, align 4
    %9 = load i32, ptr %4, align 4
    %10 = icmp slt i32 %8, %9
    br i1 %10, label %11, label %22

11:                                   ; preds = %7
    %12 = load ptr, ptr %3, align 8
    %13 = load i32, ptr %6, align 4
    %14 = sext i32 %13 to i64
    %15 = getelementptr inbounds [i32], ptr %12, i64, %14
```


32775.977f0

```
1 @btime cpp_sum($vec_float)
```



```
279.494 μs (0 allocations: 0 bytes)
```



```
1 cpp_sum(x::Vector{Cfloat}) = ccall(("c_sum", cpp_sum_float_lib), Cfloat,  
  (Ptr{Cfloat}, Cint), x, length(x));
```

```
#include <vector>
```

```
int my_sum(std::vector<int> vec) {  
    int total = 0;  
    for (int i = 0; i < vec.size(); i++) {  
        total += vec[i];  
    }  
    return total;  
}
```

```
extern "C" {  
int c_sum(int *array, int length) {  
    std::vector<int> v;  
    v.assign(array, array + length);  
    return my_sum(v);  
}}
```

No pragma ▼

No pragma ▼

No pragma ▼

Element type : int ▼

Optimization level : -O0 ▼

☐ -msse3

☐ -mavx2

☐ -mavx512f

☐ -ffast-math

☐ -march=native

⚠ PlutoRunner.embed_display(x) is deprecated and will be removed in a future release. Use AbstractPlutoDingetjes.Display.@embed(x) instead.

Tip

Easily call C++ code from Julia or Python by adding a C interface like the `c_sum` in this example.

LLVM Superword-Level Parallelism (SLP) Vectorizer ↻

f (generic function with 2 methods)

```
1 f(a, b) = (a[1] + b[1], a[2] + b[2], a[3] + b[3], a[4] + b[4])
```


```
1 @code_llvm debuginfo=:none f((1, 2, 3, 4), (5, 6, 7, 8))
```



```
; Function Signature: f(NTuple{4, Int64}, NTuple{4, Int64})
define void @julia_f_34911(ptr noalias nocapture noundef nonnull sret([4 x i64]) align 8 dereferenceable(32) %sret_return, ptr nocapture noundef nonnull readonly align 8 dereferenceable(32) %"a::Tuple", ptr nocapture noundef nonnull readonly align 8 dereferenceable(32) %"b::Tuple") #0 {
top:
    %0 = load <4 x i64>, ptr %"a::Tuple", align 8
    %1 = load <4 x i64>, ptr %"b::Tuple", align 8
    %2 = add <4 x i64> %1, %0
    store <4 x i64> %2, ptr %sret_return, align 8
    ret void
}
```



Inspection with godbolt Compiler Explorer [↗](#)



[Example source](#)

Interleaving ⇄

```
; ModuleID = '/tmp/jl_60dPrE/main.c'
source_filename = "/tmp/jl_60dPrE/main.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-f8
0:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

; Function Attrs: noinline nounwind optnone uwtable
define dso_local void @add(ptr noundef %0, i32 noundef %1) #0 {
    %3 = alloca ptr, align 8
    %4 = alloca i32, align 4
    %5 = alloca i32, align 4
    %6 = alloca i32, align 4
    store ptr %0, ptr %3, align 8
    store i32 %1, ptr %4, align 4
    store i32 0, ptr %5, align 4
    store i32 0, ptr %6, align 4
    br label %7

7:                                     ; preds = %22, %2
    %8 = load i32, ptr %6, align 4
    %9 = load i32, ptr %4, align 4
    %10 = icmp slt i32 %8, %9
    br i1 %10, label %11, label %25

11:                                   ; preds = %7
    %12 = load ptr, ptr %3, align 8
    %13 = load i32, ptr %6, align 4
    %14 = sext i32 %13 to i64
    %15 = getelementptr inbounds [i32], ptr %12, i64 %14
```

With `interleave_count(2)` and

- -O1, we see `!{"llvm.loop.interleave.count", i32 2}` at the end but it hasn't been applied
- -O2, we see the interleaving is applied

```
void add(int *a, int length) {
    int total = 0;
    for (int i = 0; i < length; i++) {
        a[i] = a[i] + 1;
    }
    return;
}
```

No pragma ▼

No pragma ▼

No pragma ▼

Element type : int ▼

Optimization level : -O0 ▼

☐ -msse3

☐ -mavx2

☐ -mavx512f

☐ -ffast-math

☐ -march=native

⚠ `PlutoRunner.embed_display(x)` is deprecated and will be removed in a future release. Use `AbstractPlutoDingetjes.Display.@embed(x)` instead.

Modern CPU architectures

Reorder buffer:

```
a = b + c;  
d = a + c; // Needs to wait for a  
e = b - c; // Independent from a  
           // Can be executed before
```


Execution ports

- Arithmetic Logic Unit (ALU)

```
a + b
```

- Address Generation Unit (AGU)
 - Load AGU in port 2 & 3
 - Store AGU in port 4

```
a[i] -> a + i * sizeof(...)
```



Why reordering ?

```
load  r1 ← a  
load  r2 ← b  
load  r3 ← c  
load  r4 ← d  
add   r5 ← r1 + r2  
add   r6 ← r3 + r4  
add   r7 ← r6 + 1
```

Assume

- L1 latency is 4 cycles
- L2 latency is 12 cycles
- a is missing from L1 cache but is in L2 cache

Without reordering

Cycle	ALU	ALU	Load	Load
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				
13				
14				
15				
16				
17				
18				
19				
20				
21				
22				
23				
24				
25				
26				
27				
28				
29				
30				
31				
32				
33				
34				
35				
36				
37				
38				
39				
40				
41				
42				
43				
44				
45				
46				
47				
48				
49				
50				
51				
52				
53				
54				
55				
56				
57				
58				
59				
60				
61				
62				
63				
64				
65				
66				
67				
68				
69				
70				
71				
72				
73				
74				
75				
76				
77				
78				
79				
80				
81				
82				
83				
84				
85				
86				
87				
88				
89				
90				
91				
92				
93				
94				
95				
96				
97				
98				
99				
100				

With reordering

Cycle	ALU	ALU	Load	Load
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				
13				
14				
15				
16				
17				
18				
19				
20				
21				
22				
23				
24				
25				
26				
27				
28				
29				
30				
31				
32				
33				
34				
35				
36				
37				
38				
39				
40				
41				
42				
43				
44				
45				
46				
47				
48				
49				
50				
51				
52				
53				
54				
55				
56				
57				
58				
59				
60				
61				
62				
63				
64				
65				
66				
67				
68				
69				
70				
71				
72				
73				
74				
75				
76				
77				
78				
79				
80				
81				
82				
83				
84				
85				
86				
87				
88				
89				
90				
91				
92				
93				
94				
95				
96				
97				
98				
99				
100				

1		a	b
2		c	d
3-5		L2 ⌘	L1 ⌘
6-12		L2 ⌘	
13	r1 + r2		r3 + r4
14	r6 + 1		

1		a	b
2		c	d
3-5		L2 ⌘	L1 ⌘
6	r3 + r4	L2 ⌘	
7	r6 + 1	L2 ⌘	
8-12		L2 ⌘	
13	r1 + r2		


Without unrolling ↪

```

loop:
load r1 ← a[i]
load r2 ← b[i]
add  r3 ← r1 + r2
store c[i] ← r3
add  i ← i + 1
cmp  i, N
jl   #loop ; Jump if Less

```

Cycle	ALU	ALU	Load	Load	Store	Jump
1			a[i]	b[i]		
2-4			L1 ⌘	L1 ⌘		
5	r1 + r2					
6	i + 1				c[i]	
7	i <? N					
8						#loop



Interleaving allowed by unrolling

```
loop2:
  load r1 ← a[i]
  load r2 ← b[i]
  load r4 ← a[i+1]
  load r5 ← b[i+1]
  add r3 ← r1 + r2
  add r6 ← r4 + r5
  store c[i] ← r3
  store c[i+1] ← r6
  add i ← i + 2
  cmp i, N-1
  jl #loop2
  cmp i, N
  jl #loop ; tail case, no unrolling
```



Cycle	ALU	ALU	Load	Load	Store	Jump
1			a[i]	b[i]		
2			a[i+1]	b[i+1]		
3-4			L1 ⌘	L1 ⌘		
5	r1 + r2					
6		r4 + r5			c[i]	
7	i + 2				c[i+1]	
8	i <? N-1					
9						#loop2

Further notes

- Can be controlled with `#pragma clang loop interleave_count(4)` but the compiler usually already does a good job choosing the `interleave_count`
- This count depends on whether it is *compute bound* or *memory bound*, see next lecture where we discuss these.
- The above examples assumes the value of `a[i]` etc... are in L1 cache. If not (aka "L1 cache miss"), the latency for loading `a[i]`, will be much longer.

Additional practical limits include:

Limit	Typical value	Applies when
-------	---------------	--------------

Load ports	2/cycle	Always – hard throughput ceiling
Load buffer (MOB)	~70-100 entries	All loads, hit or miss
Line fill buffers	~12 entries	L1 cache misses only
Reorder buffer	~200 μ ops	Window used for reordering

Fused instructions ⇄

Fused Multiply-Add (FMA) ⇄

```
; ModuleID = '/tmp/jl_YurY4M/main.c'
source_filename = "/tmp/jl_YurY4M/main.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-i128:128-f8
0:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

; Function Attrs: noinline nounwind optnone uwtable
define dso_local void @elementwise_muladd(ptr noundef %0, ptr noundef %1, ptr
noundef %2, i32 noundef %3) #0 {
    %5 = alloca ptr, align 8
    %6 = alloca ptr, align 8
    %7 = alloca ptr, align 8
    %8 = alloca i32, align 4
    %9 = alloca i32, align 4
    %10 = alloca i32, align 4
    store ptr %0, ptr %5, align 8
    store ptr %1, ptr %6, align 8
    store ptr %2, ptr %7, align 8
    store i32 %3, ptr %8, align 4
    store i32 0, ptr %9, align 4
    store i32 0, ptr %10, align 4
    br label %11

11:                                     ; preds = %33, %4
    %12 = load i32, ptr %10, align 4
    %13 = load i32, ptr %8, align 4
    %14 = icmp slt i32 %12, %13
    br i1 %14, label %15, label %36
```

```
void elementwise_muladd(int *c, int *a, int *b, int length) {
    int total = 0;
    for (int i = 0; i < length; i++) {
        c[i] += a[i] * b[i];
    }
    return;
}
```

No pragma ▼

No pragma ▼

No pragma ▼

Element type : int ▼

Optimization level :

- ☐ -msse3
- ☐ -mavx2
- ☐ -mavx512f
- ☐ -ffast-math
- ☐ -march=native

⚠ `PlutoRunner.embed_display(x)` is deprecated and will be removed in a future release. Use `AbstractPlutoDingetjes.Display.@embed(x)` instead.

Explicit FMA: ☐

Further readings

.....

Slides inspired from:

- [SIMD in Julia](#)
- [Demystifying Auto-vectorization in Julia](#)
- [Auto-Vectorization in LLVM](#)
- [Information on processor architecture](#)
- [Advanced books on optimization](#)

The End

```
1 using SimpleClang, PlutoUI, PlutoUI.ExperimentalLayout, HypertextLiteral; @luxor, @html_
   StaticArrays, BenchmarkTools, PlutoTeachingTools
```

img (generic function with 3 methods)

```
1 begin
2   struct Path
3     path::String
4   end
5
6   function imgpath(path::Path)
7     file = path.path
8     if !('.' in file)
9       file = file * ".png"
10    end
11    return joinpath(joinpath(@__DIR__, "images", file))
12  end
13
14  function img(path::Path, args...; kws...)
15    return PlutoUI.LocalResource(imgpath(path), args...)
16  end
17
18  struct URL
19    url::String
20  end
21
22  function save_image(url::URL, html_attributes...; name = split(url.url, '/') [end],
23    kws...)
24    path = joinpath("cache", name)
25    return PlutoTeachingTools.RobustLocalResource(url.url, path,
26      html_attributes...), path
27  end
28
29  function img(url::URL, args...; kws...)
30    r, _ = save_image(url, args...; kws...)
31    return @html("<a href=$(url.url)>$r</a>")
32  end
33
34  function img(file::String, args...; kws...)
35    if startswith(file, "http")
36      img(URL(file), args...; kws...)
37    else
38      img(Path(file), args...; kws...)
39    end
40  end
41 end
```

qa (generic function with 2 methods)

```
1 begin
2 function qa(question, answer)
3     return @html("<details><summary>$question</summary>$answer</details>")
4 end
5 function _inline_html(m::Markdown.Paragraph)
6     return sprint(Markdown.htmlinline, m.content)
7 end
8 function qa(question::Markdown.MD, answer)
9     # 'html(question)' will create '<p>' if 'question.content[]' is
    'Markdown.Paragraph'
10    # This will print the question on a new line and we don't want that:
11    h = HTML(_inline_html(question.content[]))
12    return qa(h, answer)
13 end
14 end
```

```
1 function Luxor.placeimage(url::URL, pos; scale = 1.0, centered = true, kws...)
2     r, path = save_image(url; kws...)
3     if r.mime isa MIME"image/svg+xml"
4         img = Luxor.readsvg(path)
5     else
6         img = Luxor.readpng(path)
7     end
8     Luxor.gsave()
9     Luxor.scale(scale)
10    Luxor.placeimage(img, pos / scale; centered)
11    Luxor.grestore() # undo 'Luxor.scale'
12    return
13 end
```

CenteredBoundingBox (generic function with 1 method)

```
1 function CenteredBoundingBox(str)
2     xbearing, ybearing, width, height, xadvance, yadvance =
3         Luxor.textextents(str)
4     lcorner = Luxor.Point(xbearing - width/2, ybearing + height/2)
5     ocorner = Luxor.Point(lcorner.x + width, lcorner.y + height)
6     return Luxor.BoundingBox(lcorner, ocorner)
7 end
```

boxed (generic function with 1 method)

```
1 function boxed(str::AbstractString, p; hue = "lightgrey")
2     Luxor.gsave()
3     Luxor.translate(p)
4     Luxor.sethue(hue)
5     Luxor.poly(CenteredBoundedBox(str) + 5, action = :stroke, close=true)
6     Luxor.sethue("black")
7     Luxor.text(str, Luxor.Point(0, 0); halign = :center, valign = :middle)
8     #settext("<span font='26'>$str</span>", halign="center", markup=true)
9     Luxor.origin()
10    Luxor.grestore() # strokecolor
11 end
```